



Decode op-mix vs model size — sharegpt (matmul share grows with model)



matmul attention ffn norm softmax rope residual kv_cache embedding