

matmul: prefill (compute-bound, high intensity) collapses to decode GEMV (memory-bound, ~1) — 4 models

